
Benhao Huang
huskydogewoof@gmail.com
<https://huskydoge.github.io/> | [Google Scholar](#)

EDUCATION

School of Electronics Information and Electrical Engineering, Shanghai Jiao Tong University Sep 2021 – Present
B.ENG. in Computer Science and Technology (IEEE Honor Class), GPA: 93.14/100.00, 4.08/4.30, (Rank 5/127)

Relevant Coursework:

- **Computer Science:** Design and Analysis of Algorithms (A+), Computer Networks (A+), Operating System (A+), Programming Languages and Compilers (A+), Natural Language Processing (A+), Database System Technology (A+), Computer Vision (A+), Principles and Methods of Program Design (A+), Program Design Practice (A+), Introduction to Data Science (A+)
- **Mathematics:** Mathematical Analysis (A+), Linear and Convex Optimization (A+), Information Theory (A+), Complex Analysis (A+), Probability and Statistics (A+)
- **Additional Academic Pursuits:** Engaged in a dual degree program in Mathematics and Applied Mathematics. Coursework included: Complex Analysis, Abstract Algebra, Linear Algebra II

TOEFL: 107, S24, R27, L29, W27, GRE: V157, Q170, AW4.0

RESEARCH INTERESTS

- Multi-modal World Model, World Model-based Reasoning and Planning, VLM / LLM Agents
- Interpretability AI, Data Influence Analysis, Dataset Synthesis and Distillation
- Acceleration of Generative Model

PUBLICATIONS

- DCA-Bench: A Benchmark for Dataset Curation Agents [[paper](#)]
Benhao Huang, Yingzhuo Yu, Jin Huang, Xingjian Zhang, Jiaqi W. Ma. **KDD DB Track Oral 2025**
- Contrasting Adversarial Perturbations: The Space of Harmless Perturbations [[paper](#)]
Lu Chen, Shaofeng Li, **Benhao Huang**, Fan Yang, Zheng Li, Jie Li, Yuan Luo. **AAAI 2025**
- Defining and Extracting Generalizable Interaction Primitives from DNNs. [[paper](#) | [code](#)]
Lu Chen, Siyu Lou, **Benhao Huang**, Quanshi Zhang. **ICLR 2024**.

SELECTED PROJECTS

PandoraV2: Towards General World Model with Natural Language Actions and Video States

In Progress (co-lead) Advisor: Prof. Zhiting Hu [[code](#)] Jun 2024 – Present

- **Diffusion Game Engine:** Built an auto-regressive Image-to-Video (I2V) model capable of simulating 2D platformer games (e.g., Mario), allowing control of both characters and environmental elements using text inputs on the fly. Proposed and implemented window-slide conditioning to support the generation of game videos lasting longer than one minute.
- **Video Diffusion Model Acceleration:** Spearheaded a sub-project focusing on optimizing video diffusion for real-time game generation, achieving generation speeds of under 1 second per round.
- **Complex Video Captioning:** Led a sub-project aimed at enhancing video captioning for complex scenarios where even state-of-the-art visual language models tend to falter, ensuring more accurate descriptions.
- **Large-Scale Training Data Pipeline:** Designed and implemented a high-efficiency processing pipeline for video training data, processing over 10 million videos simultaneously, significantly improving the overall data quality and processing speed.

DCA-Bench: A Benchmark for Dataset Curation Agents

KDD DB Track Oral 2025 (1st author) Advisor: Prof. Jiaqi Ma [[paper](#)] Jan 2023 – Present

- Identified a novel task for LLM agents – detecting dataset quality issues (e.g. cross-file discrepancy, hidden corruptions) for the purpose of automating AI training data curation – and developed the first benchmark for this task. We collected 221 instances and designed a four-level difficulty for more fine-grained analysis on agents' performance.
- Developed a LLM-based automatic evaluator for scalable evaluations, achieving a reliable and robust to self-preference or length bias demonstrated through comprehensive experiments.
- Led the project, conducting surveys, implementing code, completing experiments, and writings.

Defining and Extracting Generalizable Interaction Primitives from DNNs

ICLR 2024 Advisor: Prof. Quanshi Zhang [[paper](#) | [code](#)] Sep 2023 – Jan 2024

- Ideated a method to extract shared interactions of different DNNs trained for the same task.

-
- Conducted contrast experiments, illustrating that the extracted interactions can better reflect common knowledge shared by different DNNs. These shared interactions can be further used to interpret the output from DNNs.
 - Implemented the main experiment codes and engaged in algorithm design, deploying the GitHub repository.

RESEARCH EXPERIENCE

- Research Intern, [MAITRIX Lab](#), University of California San Diego. Apr 2024 – Present
Developed visual World Model with focus on game generation *Advisor: Prof. [Zhiting Hu](#)*
- Research Intern, Alignment Team, [Moonshot AI](#) Mar 2024 – Jun 2024
Explored prompt priorities alignment of LLMs *Advisor: [Flood Sung](#), [Yanan Zheng](#)*
- Research Intern, [TRAIS Lab](#), University of Illinois Urbana-Champaign Nov 2023 – Present
Constructed a LLM Agent benchmark for dataset issue detection *Advisor: Prof. [Jiaqi Ma](#)*
- Research Intern, [XAI Lab](#), Shanghai Jiao Tong University Apr 2023 – Jan 2024
Extracted common knowledge of different LLMs *Advisor: Prof. [Quanshi Zhang](#)*

AWARDS

- National Scholarship (Top 0.2% nationwide, ranked 1/127 in my major) 2023 - 2024
- Rui Yuan-Hong Shan Scholarship (Top 2%), SJTU 2022 - 2023
- Shao Qiu Scholarship (Top 4%), SJTU 2021 - 2022
- Meritorious Winner of Mathematical Contest in Modeling 2022

OTHER

- Student Mentor for CS2612 Programming Languages and Compilers 2023 - 2024
- Student Mentor for CS2601 Convex Optimization 2023 - 2024
- Volunteer at Shanghai Marathon 2022 - 2024
- Member of the Outreach Department, SJTU Spark Program Student Association 2021 - 2022